

Performance Comparison of Python Runtime Systems: CPython, PyPy, and Jython

Ege Demirci, Divyanshu Bhardwaj

CMPSC 263 Runtime Systems
UC Santa Barbara

Winter 2025

Outline

- 1 Introduction
- 2 Problem Statement
- 3 Experimental Setup
- 4 Findings
- 5 Discussion and Conclusion

Python: Advantages and Challenges

Advantages:

- High-level, dynamically typed language
- Elegant syntax and readability
- Widespread adoption across domains
- Developer productivity
- Expressive with minimal code

Challenges:

- Interpreted nature affects performance
- Overhead in CPU-bound applications
- Global Interpreter Lock (GIL) restricts concurrency
- Memory management overhead

Alternative Python Implementations

CPython Standard implementation, compiles to bytecode

- Reference implementation
- Direct C API integration
- Limited by GIL for multi-threading

PyPy Written in RPython with a Just-In-Time (JIT) compiler

- Dynamically translates code paths to machine code
- Meta-tracing JIT technique
- Potential speed improvements for long-running processes

Jython Translates Python code to Java bytecode

- Runs on the Java Virtual Machine (JVM)
- Access to Java libraries and ecosystem

Motivation for the Study

- Python increasingly used in performance-critical domains:
 - Data science
 - Machine learning
 - Financial computing
- Each interpreter has unique strengths and limitations
- No one-size-fits-all solution for all scenarios
- Need for systematic evaluation through benchmarks
- Goal: Provide empirical data for informed decisions

Problem Statement

- Performance variability depending on runtime environment
- Choosing optimal interpreter critical for time-sensitive applications
- Research questions:
 - How do CPython, PyPy, and Jython compare in terms of execution speed, memory usage, and overall performance?
 - In which scenarios does one interpreter clearly outperform the others?
 - CPU-bound computations
 - I/O operations
 - Memory-intensive tasks
- Goal: Provide data-driven recommendations for developers

Benchmark Suite Design

CPU-Intensive Benchmarks:

- **Fibonacci:** Recursive algorithm computation
- **Sort:** Large array sorting of random numbers
- **Prime Number:** Calculating primes up to a limit
- **N-Body:** Physics simulation of gravitational interactions

I/O-Intensive Benchmarks:

- **I/O:** File system operations (read/write)
- **JSON:** Serialization and deserialization
- **Asyncio:** Asynchronous task execution
- **Dictionary:** Dictionary operations (create, lookup, insert, delete)

Methodology

- Master automation script runs all benchmarks on each interpreter:
 - **CPython:** Using `python3`
 - **PyPy:** Using `pypy3`
 - **Jython:** Using `jython`
- Jython compatibility handled with alternative implementations
- Each test outputs JSON-formatted metrics:
 - Execution time
 - CPU time
 - Wall time
 - Memory usage
 - Iteration counts

Hardware and Software Environment

Software Environments:


- **CPython:** Version 3.10.12
- **PyPy:** Version 7.3.11 with Python 3 support
- **Jython:** Version 2.7.3

- Multiple executions per benchmark to ensure statistical validity
- T-tests confirmed low variance between runs
- Statistical significance tests performed on interpreter differences

Hardware Specifications:

- **CPU:** Intel Xeon 2nd generation (2 cores)
- **RAM:** 8 GB
- **OS:** Ubuntu 22.04.3 LTS
- **Storage:** 160 GB NVMe SSD

CPU-Intensive Benchmark Results



CPU-Intensive Performance Analysis

PyPy's advantages:

- JIT compilation optimizes frequently executed code

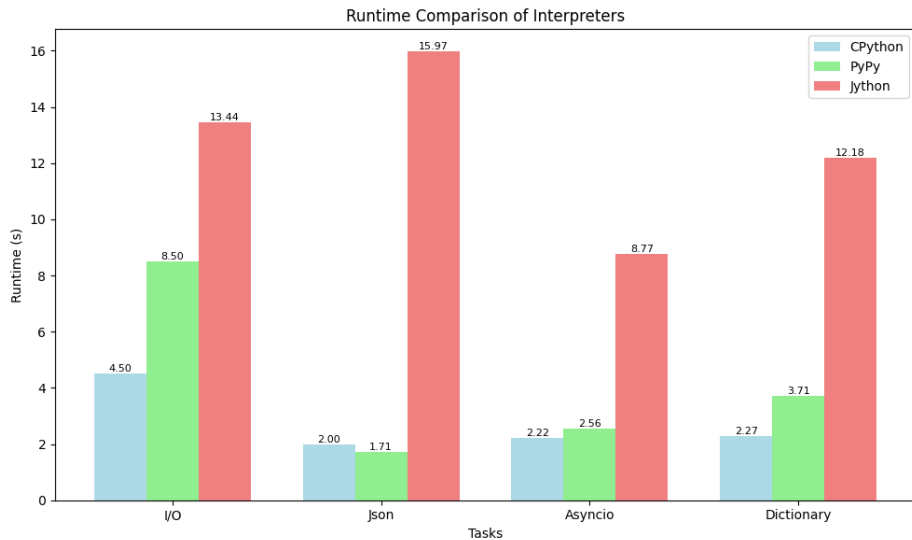
- Raw machine-level arithmetic instead of Python objects
- Effective tracing of hot paths in long-running computations

Jython's performance:

- Longer warm-up time converting Python to Java bytecode

- JVM execution introduces overhead
- Prime Number exception: JVM's optimization of arithmetic operations and loop unrolling benefits this specific workload

I/O-Intensive Benchmark Results



I/O-Intensive Performance Analysis

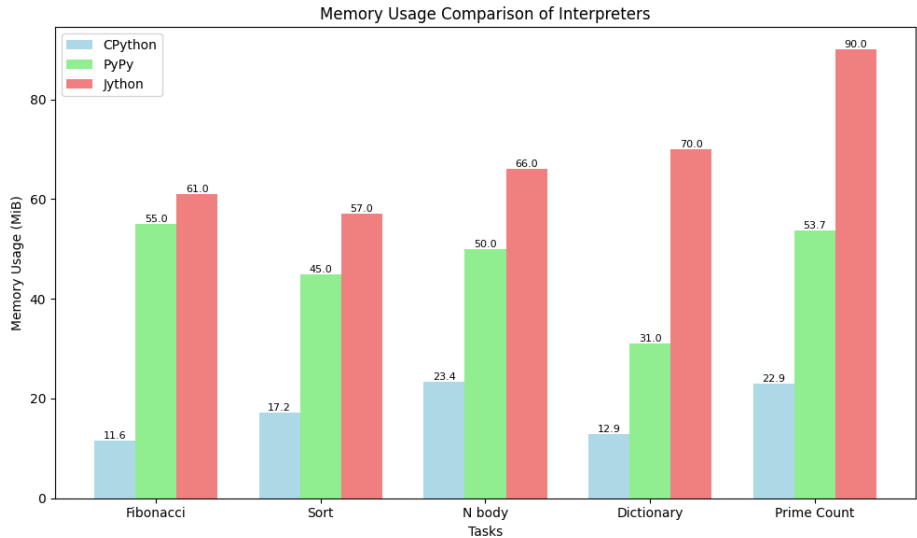
Reduced PyPy advantage: ■ I/O operations rely on system calls and external processes

- JIT compiler cannot optimize these calls
- CPython's buffering mechanisms highly optimized

Jython's challenges: ■ Overhead from Python-to-Java-bytecode conversion

- JVM execution adds another layer
- CPython's native C libraries for I/O operations more efficient

Memory Usage Comparison



Memory Usage Analysis

- CPython:
 - Efficient reference counting with cycle detection
 - Minimal memory overhead
 - Can have pauses during large object deallocation
- PyPy:
 - Generational, tracing garbage collector
 - Maintains additional data structures for JIT compiler
 - "Incminimark" incremental GC improves responsiveness but increases memory usage
- Jython:
 - Inherits JVM overhead
 - Java objects typically larger than CPython objects
 - JVM's garbage collection adds memory management overhead

Key Findings Summary

CPython

- Reliable baseline performance
- Smallest memory footprint
- Best library compatibility
- Limited by GIL for multi-threading

PyPy

- 9-10 \times faster for CPU-bound tasks
- Larger memory footprint
- JIT optimization overhead
- Limited benefit for I/O tasks

Jython

- Generally slower execution
- Highest memory usage
- Java interoperability
- No GIL limitations

No single interpreter is universally optimal

Practical Recommendations

Choose PyPy when: ■ Working with long-running, CPU-intensive applications

- Memory usage is less critical than execution speed
- Application uses pure Python code (minimal C extensions)

Choose CPython when: ■ Using libraries with native C extensions (numpy, scipy, etc.)

- Memory efficiency is important
- Compatibility with the Python ecosystem is essential
- Working with I/O-intensive applications

Choose Jython when: ■ Integration with Java libraries is required

- Working in Java-centric environments
- Need true multi-threading (no GIL)

Future Developments

- CPython improvements:
 - Potential removal of the GIL
 - JIT compilation capabilities
 - Memory optimization efforts
- PyPy enhancements:
 - Improved compatibility with C extensions
 - Memory usage optimizations
 - Faster warm-up times
- Continued evolution of Python's ecosystem warrants periodic re-evaluation of interpreter choices

Conclusion

- Performance characteristics of Python interpreters vary significantly by workload
- For CPU-intensive tasks:
 - PyPy delivers substantial speedups (9-10 \times)
- For I/O-intensive tasks:
 - CPython offers competitive or superior performance
- Memory usage considerations:
 - CPython most efficient
 - PyPy moderate
 - Jython highest
- Select interpreter based on application priorities and constraints

References

- ❶ G. van Rossum and F. L. Drake Jr., "Python 3 Reference Manual," CreateSpace, 2009.
- ❷ M. Lutz, "Learning Python: Powerful Object-Oriented Programming," 5th ed., O'Reilly Media, 2013.
- ❸ D. Beazley, "Understanding the Python GIL," in PyCon US 2010, Atlanta, GA, 2010.
- ❹ C. F. Bolz, A. Cuni, M. Fijałkowski, and A. Rigo, "Tracing the Meta-level: PyPy's Tracing JIT Compiler," ICIOOLPS, 2009.
- ❺ The Jython Project, "Jython: Python for the Java Platform," <https://www.jython.org/>
- ❻ S. Behnel et al., "Cython: The Best of Both Worlds," Computing in Science & Engineering, vol. 13, no. 2, 2011.
- ❼ A. Rigo and M. Fijałkowski, "Incremental Garbage Collector in PyPy," 2013.